

Logistic Regression

CS114 Lab 4

Kenneth Lai

February 7, 2020

Generative and Discriminative Models

- ▶ Generative models

Generative and Discriminative Models

- ▶ Generative models
 - ▶ Model joint distribution $P(c, d)$

Generative and Discriminative Models

- ▶ Generative models
 - ▶ Model joint distribution $P(c, d)$
 - ▶ $P(c|d) \propto P(d|c)P(c)$
 $\propto P(c, d)$

Generative and Discriminative Models

- ▶ Generative models
 - ▶ Model joint distribution $P(c, d)$
 - ▶ $P(c|d) \propto P(d|c)P(c)$
 $\propto P(c, d)$
 - ▶ Naïve Bayes, Hidden Markov Models, Bayesian Networks, etc.

Generative and Discriminative Models

- ▶ Generative models
 - ▶ Model joint distribution $P(c, d)$
 - ▶ $P(c|d) \propto P(d|c)P(c)$
 $\propto P(c, d)$
 - ▶ Naïve Bayes, Hidden Markov Models, Bayesian Networks, etc.
- ▶ Discriminative models

Generative and Discriminative Models

- ▶ Generative models
 - ▶ Model joint distribution $P(c, d)$
 - ▶ $P(c|d) \propto P(d|c)P(c)$
 $\propto P(c, d)$
 - ▶ Naïve Bayes, Hidden Markov Models, Bayesian Networks, etc.
- ▶ Discriminative models
 - ▶ Model conditional distribution $P(c|d)$ directly

Generative and Discriminative Models

- ▶ Generative models
 - ▶ Model joint distribution $P(c, d)$
 - ▶ $P(c|d) \propto P(d|c)P(c)$
 $\propto P(c, d)$
 - ▶ Naïve Bayes, Hidden Markov Models, Bayesian Networks, etc.
- ▶ Discriminative models
 - ▶ Model conditional distribution $P(c|d)$ directly
 - ▶ Logistic Regression, Conditional Random Fields, Neural Networks, etc.

Logistic Regression

- ▶ Suppose we observe a document $d = \text{"Chinese Chinese Chinese Tokyo Japan"}$. Is the document Chinese or Japanese?

Logistic Regression

► Training data:

document	class
Chinese Beijing Chinese	Chinese
Chinese Chinese Shanghai	Chinese
Chinese Macao	Chinese
Tokyo Japan Chinese	Japanese

Logistic Regression

- ▶ Documents are characterized by features

Logistic Regression

- ▶ Documents are characterized by features
 - ▶ No independence assumptions

Logistic Regression

- ▶ Documents are characterized by features
 - ▶ No independence assumptions
- ▶ For each feature i :
 - ▶ Value x_i
 - ▶ Weight w_i

Logistic Regression

- ▶ Documents are characterized by features
 - ▶ No independence assumptions
- ▶ For each feature i :
 - ▶ Value x_i
 - ▶ Weight w_i
- ▶ Bias term b

Logistic Regression

- ▶ Documents are characterized by features
 - ▶ No independence assumptions
- ▶ For each feature i :
 - ▶ Value x_i
 - ▶ Weight w_i
- ▶ Bias term b

- ▶ “Score” (log-odds) $z = \left(\sum_{i=1}^n w_i x_i \right) + b = \mathbf{w} \cdot \mathbf{x} + b$

Logistic Regression

- ▶ Logistic function $\sigma(z) = \frac{1}{1 + e^{-z}}$

Logistic Regression

- ▶ Logistic function $\sigma(z) = \frac{1}{1 + e^{-z}}$
- ▶ $p(y = 1|\mathbf{x}) = \sigma(\mathbf{w} \cdot \mathbf{x} + b) = \frac{1}{1 + e^{-(\mathbf{w} \cdot \mathbf{x} + b)}}$
- ▶ $p(y = 0|\mathbf{x}) = 1 - \sigma(\mathbf{w} \cdot \mathbf{x} + b) = \frac{e^{-(\mathbf{w} \cdot \mathbf{x} + b)}}{1 + e^{-(\mathbf{w} \cdot \mathbf{x} + b)}}$

Cross-entropy Loss

- ▶ Let $\hat{y} = \sigma(\mathbf{w} \cdot \mathbf{x} + b)$ be the classifier output for some \mathbf{x}

Cross-entropy Loss

- ▶ Let $\hat{y} = \sigma(\mathbf{w} \cdot \mathbf{x} + b)$ be the classifier output for some \mathbf{x}
- ▶ What is the probability that the classifier is correct?

Cross-entropy Loss

- ▶ Let $\hat{y} = \sigma(\mathbf{w} \cdot \mathbf{x} + b)$ be the classifier output for some \mathbf{x}
- ▶ What is the probability that the classifier is correct?
 - ▶ If $y = 1$, then $P(y = 1|\mathbf{x}) = \hat{y}$
 - ▶ If $y = 0$, then $P(y = 0|\mathbf{x}) = 1 - \hat{y}$

Cross-entropy Loss

- ▶ In general, $P(y|\mathbf{x}) = \hat{y}^y(1 - \hat{y})^{1-y}$

Cross-entropy Loss

- ▶ In general, $P(y|\mathbf{x}) = \hat{y}^y(1 - \hat{y})^{1-y}$
- ▶ Take the log of both sides:
$$\log P(y|\mathbf{x}) = y \log \hat{y} + (1 - y) \log(1 - \hat{y})$$

Cross-entropy Loss

- ▶ In general, $P(y|\mathbf{x}) = \hat{y}^y(1 - \hat{y})^{1-y}$
- ▶ Take the log of both sides:
$$\log P(y|\mathbf{x}) = y \log \hat{y} + (1 - y) \log(1 - \hat{y})$$
- ▶ Turn this into a loss function:
$$L(\hat{y}, y) = -\log P(y|\mathbf{x}) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$$

Cross-entropy Loss

- ▶ Minimize average loss for each example j :

$$Cost(\mathbf{w}, b) = \frac{1}{m} \sum_{j=1}^m L(\hat{y}^{(j)}, y^{(j)})$$

Cross-entropy Loss

- ▶ Minimize average loss for each example j :

$$Cost(\mathbf{w}, b) = \frac{1}{m} \sum_{j=1}^m L(\hat{y}^{(j)}, y^{(j)})$$

- ▶ How?

Gradient Descent




Figure 5.3 The first step in iteratively finding the minimum of this loss function, by moving w in the reverse direction from the slope of the function. Since the slope is negative, we need to move w in a positive direction, to the right. Here superscripts are used for learning steps, so w^1 means the initial value of w (which is 0), w^2 at the second step, and so on.

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)
- ▶ At each time step t :

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ Compute gradient ∇L

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ Compute gradient ∇L

$$\nabla L = \begin{bmatrix} \frac{\partial L}{\partial w_1} \\ \vdots \\ \frac{\partial L}{\partial w_n} \\ \frac{\partial L}{\partial b} \end{bmatrix}$$

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ Compute gradient ∇L

- ▶ $\nabla L = \begin{bmatrix} \frac{\partial L}{\partial w_1} \\ \vdots \\ \frac{\partial L}{\partial w_n} \\ \frac{\partial L}{\partial b} \end{bmatrix}$
 - ▶ \approx slope of loss function L

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ randomly
- ▶ At each time step t :
 - ▶ Compute gradient ∇L
 - ▶ Move in direction of negative gradient

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ randomly
- ▶ At each time step t :
 - ▶ Compute gradient ∇L
 - ▶ Move in direction of negative gradient
- ▶ $\theta_{t+1} = \theta_t - \eta \nabla L$

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ randomly
- ▶ At each time step t :
 - ▶ Compute gradient ∇L
 - ▶ Move in direction of negative gradient
- ▶ $\theta_{t+1} = \theta_t - \eta \nabla L$
 - ▶ η = learning rate

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ randomly
- ▶ At each time step t :
 - ▶ Compute gradient ∇L
 - ▶ Move in direction of negative gradient
- ▶ $\theta_{t+1} = \theta_t - \eta \nabla L$
 - ▶ η = learning rate
 - ▶ “Hyperparameter”: parameter set before training

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ randomly
- ▶ At each time step t :
 - ▶ Compute gradient ∇L
 - ▶ Move in direction of negative gradient
- ▶ $\theta_{t+1} = \theta_t - \eta \nabla L$
 - ▶ η = learning rate
 - ▶ “Hyperparameter”: parameter set before training
 - ▶ Trade-off between speed of convergence and “zig-zag” behavior

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ randomly
- ▶ At each time step t :
 - ▶ Compute gradient ∇L
 - ▶ Move in direction of negative gradient
- ▶ $\theta_{t+1} = \theta_t - \eta \nabla L$
 - ▶ η = learning rate
 - ▶ “Hyperparameter”: parameter set before training
 - ▶ Trade-off between speed of convergence and “zig-zag” behavior
 - ▶ Often a function of t

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ randomly
- ▶ At each time step t :
 - ▶ Compute gradient ∇L
 - ▶ Move in direction of negative gradient
- ▶ $\theta_{t+1} = \theta_t - \eta \nabla L$
 - ▶ η = learning rate
 - ▶ “Hyperparameter”: parameter set before training
 - ▶ Trade-off between speed of convergence and “zig-zag” behavior
 - ▶ Often a function of t
- ▶ Because L is convex, we eventually reach a global minimum

Gradient Descent

► $L(\mathbf{w}, b) = -[y \log \sigma(\mathbf{w} \cdot \mathbf{x} + b) + (1 - y) \log(1 - \sigma(\mathbf{w} \cdot \mathbf{x} + b))]$

Gradient Descent

- ▶ $L(\mathbf{w}, b) = -[y \log \sigma(\mathbf{w} \cdot \mathbf{x} + b) + (1 - y) \log(1 - \sigma(\mathbf{w} \cdot \mathbf{x} + b))]$
- ▶ ...
- ▶ (calculus)
- ▶ ...

Gradient Descent

- ▶ $L(\mathbf{w}, b) = -[y \log \sigma(\mathbf{w} \cdot \mathbf{x} + b) + (1 - y) \log(1 - \sigma(\mathbf{w} \cdot \mathbf{x} + b))]$
- ▶ ...
- ▶ (calculus)
- ▶ ...
- ▶ $\frac{\partial L}{\partial w_i} = [\sigma(\mathbf{w} \cdot \mathbf{x} + b) - y]x_i$
- ▶ $\frac{\partial L}{\partial b} = \sigma(\mathbf{w} \cdot \mathbf{x} + b) - y$

Example

- ▶ Python time!

Gradient Descent

- ▶ Stochastic gradient descent: update θ after every training example

Gradient Descent

- ▶ Stochastic gradient descent: update θ after every training example
 - ▶ Can result in very choppy movements

Gradient Descent

- ▶ Stochastic gradient descent: update θ after every training example
 - ▶ Can result in very choppy movements
- ▶ Batch gradient descent: update θ after processing the entire training set

Gradient Descent

- ▶ Stochastic gradient descent: update θ after every training example
 - ▶ Can result in very choppy movements
- ▶ Batch gradient descent: update θ after processing the entire training set
- ▶ Minibatch gradient descent: update θ after m training examples

Gradient Descent

- ▶ Stochastic gradient descent: update θ after every training example
 - ▶ Can result in very choppy movements
- ▶ Batch gradient descent: update θ after processing the entire training set
- ▶ Minibatch gradient descent: update θ after m training examples
 - ▶ Gradient = sum (Eisenstein) or average (Jurafsky and Martin) of individual gradients—use average for HW

Gradient Descent

- ▶ Stochastic gradient descent: update θ after every training example
 - ▶ Can result in very choppy movements
- ▶ Batch gradient descent: update θ after processing the entire training set
- ▶ Minibatch gradient descent: update θ after m training examples
 - ▶ Gradient = sum (Eisenstein) or average (Jurafsky and Martin) of individual gradients—use average for HW

$$\frac{\partial \text{Cost}}{\partial w_i} = \frac{1}{m} \sum_{j=1}^m [\sigma(\mathbf{w} \cdot \mathbf{x}^{(j)} + b) - y^{(j)}] x_i^{(j)}$$

Regularization

- ▶ Problem of overfitting

Regularization

- ▶ Problem of overfitting
 - ▶ Models can fit the training data too well

Regularization

- ▶ Problem of overfitting
 - ▶ Models can fit the training data too well
 - ▶ Accidental correlations get high weights

Regularization

- ▶ Problem of overfitting
 - ▶ Models can fit the training data too well
 - ▶ Accidental correlations get high weights
 - ▶ Poor generalization performance

Regularization

$$Cost(\mathbf{w}, b) = \left(\frac{1}{m} \sum_{j=1}^m L(\hat{y}^{(j)}, y^{(j)}) \right) + \alpha R(\mathbf{w})$$

Regularization

$$\text{Cost}(\mathbf{w}, b) = \left(\frac{1}{m} \sum_{j=1}^m L(\hat{y}^{(j)}, y^{(j)}) \right) + \alpha R(\mathbf{w})$$

- ▶ $R(\mathbf{w})$ = regularization term
- ▶ α = amount of regularization

Regularization

$$\text{Cost}(\mathbf{w}, b) = \left(\frac{1}{m} \sum_{j=1}^m L(\hat{y}^{(j)}, y^{(j)}) \right) + \alpha R(\mathbf{w})$$

- ▶ $R(\mathbf{w})$ = regularization term
- ▶ α = amount of regularization
 - ▶ Another hyperparameter

L1 Regularization

$$R(\mathbf{w}) = \|\mathbf{w}\|_1 = \sum_{i=1}^n |w_i|$$

L1 Regularization

$$R(\mathbf{w}) = \|\mathbf{w}\|_1 = \sum_{i=1}^n |w_i|$$

- ▶ = sum of absolute values of weights
- ▶ Manhattan distance
- ▶ Lasso regression
- ▶ Some large weights, many zero weights

L2 Regularization

$$R(\mathbf{w}) = \|\mathbf{w}\|_2^2 = \sum_{i=1}^n w_i^2$$

L2 Regularization

$$R(\mathbf{w}) = \|\mathbf{w}\|_2^2 = \sum_{i=1}^n w_i^2$$

- ▶ = sum of squares of weights
- ▶ Euclidean distance
- ▶ Ridge regression
- ▶ Many small weights

Multinomial Logistic Regression

- ▶ (aka maximum entropy classifier)

Multinomial Logistic Regression

- ▶ (aka maximum entropy classifier)
- ▶ Logistic regression with more than two classes

Multinomial Logistic Regression

- ▶ Separate weights and bias terms for each class c

Multinomial Logistic Regression

- ▶ Separate weights and bias terms for each class c

$$z_c = \left(\sum_{i=1}^n w_{i,c} x_i \right) + b_c = \mathbf{w}_c \cdot \mathbf{x} + b_c$$

Multinomial Logistic Regression

- ▶ Separate weights and bias terms for each class c

$$z_c = \left(\sum_{i=1}^n w_{i,c} x_i \right) + b_c = \mathbf{w}_c \cdot \mathbf{x} + b_c$$

- ▶ Softmax function $\sigma(z_c) = \frac{e^{z_c}}{\sum_{k \in C} e^{z_k}} = P(y = c | \mathbf{x})$