

From Logistic Regression to Neural Networks

CS114 Lab 5

Kenneth Lai

February 14, 2020

Logistic Regression

- ▶ For each feature i :
 - ▶ Value x_i
 - ▶ Weight w_i

Logistic Regression

- ▶ For each feature i :
 - ▶ Value x_i
 - ▶ Weight w_i
- ▶ Bias term b

Logistic Regression

- ▶ For each feature i :
 - ▶ Value x_i
 - ▶ Weight w_i
- ▶ Bias term b

- ▶ “Score” (log-odds) $z = \left(\sum_{i=1}^n w_i x_i \right) + b = \mathbf{w} \cdot \mathbf{x} + b$

Logistic Regression

- ▶ For each feature i :
 - ▶ Value x_i
 - ▶ Weight w_i
- ▶ Bias term b

- ▶ “Score” (log-odds) $z = \left(\sum_{i=1}^n w_i x_i \right) + b = \mathbf{w} \cdot \mathbf{x} + b$

- ▶ Logistic function $\sigma(z) = \frac{1}{1 + e^{-z}} = \hat{y}$

Logistic Regression

- ▶ For each feature i :

- ▶ Value x_i
- ▶ Weight w_i

- ▶ Bias term b

- ▶ “Score” (log-odds) $z = \left(\sum_{i=1}^n w_i x_i \right) + b = \mathbf{w} \cdot \mathbf{x} + b$

- ▶ Logistic function $\sigma(z) = \frac{1}{1 + e^{-z}} = \hat{y}$

- ▶ Cross-entropy loss $L = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$

Gradient Descent

- Compute gradient $\nabla L = \begin{bmatrix} \frac{\partial L}{\partial w_1} \\ \vdots \\ \frac{\partial L}{\partial w_n} \\ \frac{\partial L}{\partial b} \end{bmatrix}$, where

Gradient Descent

► Compute gradient $\nabla L = \begin{bmatrix} \frac{\partial L}{\partial w_1} \\ \vdots \\ \frac{\partial L}{\partial w_n} \\ \frac{\partial L}{\partial b} \end{bmatrix}$, where

- $\frac{\partial L}{\partial w_i} = (\hat{y} - y)x_i$
- $\frac{\partial L}{\partial b} = \hat{y} - y$

Multinomial Logistic Regression

- ▶ Separate weights and bias terms for each class c

Multinomial Logistic Regression

- ▶ Separate weights and bias terms for each class c

$$z_c = \left(\sum_{i=1}^n w_{ic} x_i \right) + b_c = \mathbf{w}_c \cdot \mathbf{x} + b_c$$

Multinomial Logistic Regression

- ▶ Separate weights and bias terms for each class c

$$z_c = \left(\sum_{i=1}^n w_{ic} x_i \right) + b_c = \mathbf{w}_c \cdot \mathbf{x} + b_c$$

- ▶ Softmax function $\sigma(z_c) = \frac{e^{z_c}}{\sum_{k \in C} e^{z_k}} = P(y = c | \mathbf{x}) = \hat{y}_c$

Multinomial Logistic Regression

- ▶ Separate weights and bias terms for each class c

$$z_c = \left(\sum_{i=1}^n w_{ic} x_i \right) + b_c = \mathbf{w}_c \cdot \mathbf{x} + b_c$$

- ▶ Softmax function $\sigma(z_c) = \frac{e^{z_c}}{\sum_{k \in C} e^{z_k}} = P(y = c | \mathbf{x}) = \hat{y}_c$
- ▶ Cross-entropy loss $L = - \sum_{k \in C} y_k \log \hat{y}_k$

Multinomial Logistic Regression

- ▶ Gradient ∇L becomes a matrix, where

Multinomial Logistic Regression

- ▶ Gradient ∇L becomes a matrix, where


- ▶ $\frac{\partial L}{\partial w_{ic}} = (\hat{y}_c - y_c)x_i$

- ▶ $\frac{\partial L}{\partial b_c} = \hat{y}_c - y_c$


Example

- ▶ Python time!


Graphical Representation of Logistic Regression



Graphical Representation of a Neuron




Graphical Representation of a Neuron




- Inputs (to neuron j)

Graphical Representation of a Neuron




- Weights (and bias term)

Graphical Representation of a Neuron




- Input function (almost always Σ)

Graphical Representation of a Neuron




► “Score”

Graphical Representation of a Neuron



- Activation function (logistic, softmax, tanh, ReLU, etc.)

Graphical Representation of a Neuron



- Activation (output of neuron j)

Gradients in Logistic Regression

- ▶ We want to compute $\frac{\partial L}{\partial w_i}$

Gradients in Logistic Regression

- ▶ We want to compute $\frac{\partial L}{\partial w_i}$
- ▶ Chain Rule of calculus: $\frac{dy}{dx} = \frac{dy}{dz} \frac{dz}{dx}$

Gradients in Logistic Regression

- ▶ We want to compute $\frac{\partial L}{\partial w_i}$
- ▶ Chain Rule of calculus: $\frac{dy}{dx} = \frac{dy}{dz} \frac{dz}{dx}$
- ▶ Looking at the graph: $\frac{\partial L}{\partial w_i} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z} \frac{\partial z}{\partial w_i}$

Gradients in Logistic Regression

$$\blacktriangleright \frac{\partial L}{\partial w_i} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z} \frac{\partial z}{\partial w_i}$$

Gradients in Logistic Regression

$$\begin{aligned} \blacktriangleright \frac{\partial L}{\partial w_i} &= \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z} x_i \\ \blacktriangleright \frac{\partial z}{\partial w_i} &= x_i \end{aligned}$$

Gradients in Logistic Regression

- ▶ $\frac{\partial L}{\partial w_i} = \frac{\partial L}{\partial \hat{y}} \hat{y}(1 - \hat{y}) x_i$
 - ▶ $\frac{\partial z}{\partial w_i} = x_i$
 - ▶ For the logistic function: $\frac{\partial \hat{y}}{\partial z} = \hat{y}(1 - \hat{y})$

Gradients in Logistic Regression

- ▶ $\frac{\partial L}{\partial w_i} = \frac{\hat{y} - y}{\hat{y}(1 - \hat{y})} \hat{y}(1 - \hat{y}) x_i$
 - ▶ $\frac{\partial z}{\partial w_i} = x_i$
 - ▶ For the logistic function: $\frac{\partial \hat{y}}{\partial z} = \hat{y}(1 - \hat{y})$
 - ▶ For cross-entropy loss: $\frac{\partial L}{\partial \hat{y}} = \frac{\hat{y} - y}{\hat{y}(1 - \hat{y})}$

Gradients in Logistic Regression

- ▶ $\frac{\partial L}{\partial w_i} = \frac{\hat{y} - y}{\hat{y}(1 - \hat{y})} \hat{y}(1 - \hat{y}) x_i = (\hat{y} - y) x_i$
 - ▶ $\frac{\partial z}{\partial w_i} = x_i$
 - ▶ For the logistic function: $\frac{\partial \hat{y}}{\partial z} = \hat{y}(1 - \hat{y})$
 - ▶ For cross-entropy loss: $\frac{\partial L}{\partial \hat{y}} = \frac{\hat{y} - y}{\hat{y}(1 - \hat{y})}$

Gradients in Feedforward Neural Networks

- For an output neuron:

Gradients in Feedforward Neural Networks

- ▶ For an output neuron:

- ▶ $\frac{\partial L}{\partial w_{ij}} = \frac{a_j - y}{a_j(1 - a_j)} a_j(1 - a_j) x_{ij} = (a_j - y) x_{ij}$

- ▶ $\frac{\partial z_j}{\partial w_{ij}} = x_{ij}$

- ▶ For the logistic function: $\frac{\partial a_j}{\partial z_j} = a_j(1 - a_j)$

- ▶ For cross-entropy loss: $\frac{\partial L}{\partial a_j} = \frac{a_j - y}{a_j(1 - a_j)}$

Gradients in Feedforward Neural Networks

- ▶ For an output neuron:

- ▶ $\frac{\partial L}{\partial w_{ij}} = \frac{a_j - y}{a_j(1 - a_j)} a_j(1 - a_j) x_{ij} = (a_j - y) x_{ij}$

- ▶ $\frac{\partial z_j}{\partial w_{ij}} = x_{ij}$

- ▶ For the logistic function: $\frac{\partial a_j}{\partial z_j} = a_j(1 - a_j)$

- ▶ For cross-entropy loss: $\frac{\partial L}{\partial a_j} = \frac{a_j - y}{a_j(1 - a_j)}$

- ▶ Note that for an output neuron, $a_j = \hat{y}$

Gradients in Feedforward Neural Networks

- For a non-output neuron:

Gradients in Feedforward Neural Networks

- ▶ For a non-output neuron:

- ▶
$$\frac{\partial L}{\partial w_{ij}} = \frac{\partial L}{\partial a_j} \frac{\partial a_j}{\partial z_j} \frac{\partial z_j}{\partial w_{ij}}$$

Gradients in Feedforward Neural Networks

- ▶ For a non-output neuron:

- ▶ $\frac{\partial L}{\partial w_{ij}} = \frac{\partial L}{\partial a_j} a_j(1 - a_j)x_{ij}$

- ▶ $\frac{\partial z_j}{\partial w_{ij}} = x_{ij}$

- ▶ For the logistic function: $\frac{\partial a_j}{\partial z_j} = a_j(1 - a_j)$

Gradients in Feedforward Neural Networks

- ▶ For a non-output neuron:

- ▶ $\frac{\partial L}{\partial w_{ij}} = \frac{\partial L}{\partial a_j} a_j(1 - a_j)x_{ij}$

- ▶ $\frac{\partial z_j}{\partial w_{ij}} = x_{ij}$

- ▶ For the logistic function: $\frac{\partial a_j}{\partial z_j} = a_j(1 - a_j)$

- ▶ $\frac{\partial L}{\partial a_j} = ?$

- ▶ Note that for a non-output neuron, $a_j \neq \hat{y}$

Gradients in Feedforward Neural Networks

- ▶ Note that for a non-output neuron, $a_j = x_{j\ell}$ for some other neuron(s) ℓ

Gradients in Feedforward Neural Networks

- ▶ Note that for a non-output neuron, $a_j = x_{j\ell}$ for some other neuron(s) ℓ
 - ▶ Let \mathcal{L} be the set of all such ℓ

Gradients in Feedforward Neural Networks

- ▶ Note that for a non-output neuron, $a_j = x_{j\ell}$ for some other neuron(s) ℓ
 - ▶ Let \mathcal{L} be the set of all such ℓ
- ▶ Chain Rule of multivariable calculus:

$$\frac{df(g_1(x), \dots, g_n(x))}{dx} = \sum_{i=1}^n \frac{\partial f}{\partial g_i(x)} \frac{dg_i(x)}{dx}$$

Gradients in Feedforward Neural Networks

- ▶ Note that for a non-output neuron, $a_j = x_{j\ell}$ for some other neuron(s) ℓ
 - ▶ Let \mathcal{L} be the set of all such ℓ
- ▶ Chain Rule of multivariable calculus:

$$\frac{df(g_1(x), \dots, g_n(x))}{dx} = \sum_{i=1}^n \frac{\partial f}{\partial g_i(x)} \frac{dg_i(x)}{dx}$$

- ▶ Express L as a function of z_ℓ : $\frac{\partial L}{\partial a_j} = \sum_{\ell \in \mathcal{L}} \frac{\partial L}{\partial z_\ell} \frac{\partial z_\ell}{\partial x_{j\ell}}$

Gradients in Feedforward Neural Networks

- ▶ For a non-output neuron:

- ▶ $\frac{\partial L}{\partial w_{ij}} = \frac{\partial L}{\partial a_j} a_j(1 - a_j)x_{ij}$

- ▶ $\frac{\partial z_j}{\partial w_{ij}} = x_{ij}$

- ▶ For the logistic function: $\frac{\partial a_j}{\partial z_j} = a_j(1 - a_j)$

Gradients in Feedforward Neural Networks

- ▶ For a non-output neuron:

- ▶ $\frac{\partial L}{\partial w_{ij}} = \left(\sum_{\ell \in \mathcal{L}} \frac{\partial L}{\partial z_{\ell}} \frac{\partial z_{\ell}}{\partial x_{j\ell}} \right) a_j(1 - a_j)x_{ij}$

- ▶ $\frac{\partial z_j}{\partial w_{ij}} = x_{ij}$

- ▶ For the logistic function: $\frac{\partial a_j}{\partial z_j} = a_j(1 - a_j)$

- ▶ $\frac{\partial L}{\partial a_j} = \sum_{\ell \in \mathcal{L}} \frac{\partial L}{\partial z_{\ell}} \frac{\partial z_{\ell}}{\partial x_{j\ell}}$

Gradients in Feedforward Neural Networks

- ▶ For a non-output neuron:

- ▶ $\frac{\partial L}{\partial w_{ij}} = \left(\sum_{\ell \in \mathcal{L}} \frac{\partial L}{\partial z_{\ell}} w_{j\ell} \right) a_j(1 - a_j) x_{ij}$

- ▶ $\frac{\partial z_j}{\partial w_{ij}} = x_{ij}$

- ▶ For the logistic function: $\frac{\partial a_j}{\partial z_j} = a_j(1 - a_j)$

- ▶ $\frac{\partial L}{\partial a_j} = \sum_{\ell \in \mathcal{L}} \frac{\partial L}{\partial z_{\ell}} \frac{\partial z_{\ell}}{\partial x_{j\ell}}$

- ▶ $\frac{\partial z_{\ell}}{\partial x_{j\ell}} = w_{j\ell}$

Gradients in Feedforward Neural Networks

- ▶ For a non-output neuron:

- ▶ $\frac{\partial L}{\partial w_{ij}} = \left(\sum_{\ell \in \mathcal{L}} \frac{\partial L}{\partial z_{\ell}} w_{j\ell} \right) a_j(1 - a_j) x_{ij}$

- ▶ $\frac{\partial z_j}{\partial w_{ij}} = x_{ij}$

- ▶ For the logistic function: $\frac{\partial a_j}{\partial z_j} = a_j(1 - a_j)$

- ▶ $\frac{\partial L}{\partial a_j} = \sum_{\ell \in \mathcal{L}} \frac{\partial L}{\partial z_{\ell}} \frac{\partial z_{\ell}}{\partial x_{j\ell}}$

- ▶ $\frac{\partial z_{\ell}}{\partial x_{j\ell}} = w_{j\ell}$

- ▶ What about $\frac{\partial L}{\partial z_{\ell}}$?

Backpropagation

- ▶ We can compute $\frac{\partial L}{\partial w_{j\ell}} = \frac{\partial L}{\partial a_\ell} \frac{\partial a_\ell}{\partial z_\ell} \frac{\partial z_\ell}{\partial w_{j\ell}}$ for an output neuron ℓ

Backpropagation

- ▶ We can compute $\frac{\partial L}{\partial w_{j\ell}} = \frac{\partial L}{\partial a_\ell} \frac{\partial a_\ell}{\partial z_\ell} \frac{\partial z_\ell}{\partial w_{j\ell}}$ for an output neuron ℓ
- ▶ If we have already computed $\frac{\partial L}{\partial w_{j\ell}}$ for some neuron ℓ , then we have also computed $\frac{\partial L}{\partial z_\ell} \left(= \frac{\partial L}{\partial a_\ell} \frac{\partial a_\ell}{\partial z_\ell} \right)$

Backpropagation

- ▶ We can compute $\frac{\partial L}{\partial w_{j\ell}} = \frac{\partial L}{\partial a_\ell} \frac{\partial a_\ell}{\partial z_\ell} \frac{\partial z_\ell}{\partial w_{j\ell}}$ for an output neuron ℓ
- ▶ If we have already computed $\frac{\partial L}{\partial w_{j\ell}}$ for some neuron ℓ , then we have also computed $\frac{\partial L}{\partial z_\ell} \left(= \frac{\partial L}{\partial a_\ell} \frac{\partial a_\ell}{\partial z_\ell} \right)$
- ▶ We can then use $\frac{\partial L}{\partial z_\ell}$ to calculate $\frac{\partial L}{\partial w_{ij}} = \left(\sum_{\ell \in \mathcal{L}} \frac{\partial L}{\partial z_\ell} w_{j\ell} \right) \frac{\partial a_j}{\partial z_j} \frac{\partial z_j}{\partial w_{ij}}$ for the previous neuron(s) j

Backpropagation

- ▶ Simplifying assumptions:

Backpropagation

- ▶ Simplifying assumptions:
 - ▶ Suppose that our neurons are grouped into a sequence of **layers**
 - ▶ Also suppose that these layers are **fully connected** (every neuron in one layer is connected to every neuron in the next layer, and no others)

Backpropagation

- ▶ Simplifying assumptions:
 - ▶ Suppose that our neurons are grouped into a sequence of **layers**
 - ▶ Also suppose that these layers are **fully connected** (every neuron in one layer is connected to every neuron in the next layer, and no others)
- ▶ Then we can use matrix multiplication

Backpropagation

- ▶ Notation:

Backpropagation

- ▶ Notation:

- ▶ Let δ_l be the vector of $\frac{\partial L}{\partial z_i}$ for all neurons i in layer l
- ▶ Let $\sigma'(\mathbf{z}_l)$ be the vector of $\frac{\partial a_i}{\partial z_i}$ for $i \in l$
- ▶ Let \odot denote the element-wise (Hadamard) product

Backpropagation

- ▶ Notation:

- ▶ Let δ_l be the vector of $\frac{\partial L}{\partial z_i}$ for all neurons i in layer l
- ▶ Let $\sigma'(\mathbf{z}_l)$ be the vector of $\frac{\partial a_i}{\partial z_i}$ for $i \in l$
- ▶ Let \odot denote the element-wise (Hadamard) product

- ▶ For all layers l :

- ▶ $\nabla_l L = \delta_l \odot \mathbf{x}_l$

Backpropagation

- ▶ Notation:
 - ▶ Let δ_l be the vector of $\frac{\partial L}{\partial z_i}$ for all neurons i in layer l
 - ▶ Let $\sigma'(\mathbf{z}_l)$ be the vector of $\frac{\partial a_i}{\partial z_i}$ for $i \in l$
 - ▶ Let \odot denote the element-wise (Hadamard) product
- ▶ For all layers l :
 - ▶ $\nabla_l L = \delta_l \odot \mathbf{x}_l$
- ▶ For an output layer \mathcal{L} :
 - ▶ $\delta_{\mathcal{L}} = \hat{\mathbf{y}} - \mathbf{y}$

Backpropagation

- ▶ Notation:
 - ▶ Let δ_I be the vector of $\frac{\partial L}{\partial z_i}$ for all neurons i in layer I
 - ▶ Let $\sigma'(\mathbf{z}_I)$ be the vector of $\frac{\partial a_i}{\partial z_i}$ for $i \in I$
 - ▶ Let \odot denote the element-wise (Hadamard) product
- ▶ For all layers I :
 - ▶ $\nabla_I L = \delta_I \odot \mathbf{x}_I$
- ▶ For an output layer \mathcal{L} :
 - ▶ $\delta_{\mathcal{L}} = \hat{\mathbf{y}} - \mathbf{y}$
- ▶ For a non-output layer J (with next layer K):
 - ▶ $\delta_J = (\mathbf{W}_K^T \cdot \delta_K) \odot \sigma'(\mathbf{z}_J)$

Backpropagation

- ▶ For each (layer of) neuron(s) j :
 - ▶ Initialize parameters $\theta_j = \mathbf{w}_j, b_j$ randomly (note: not **0**)

Backpropagation

- ▶ For each (layer of) neuron(s) j :
 - ▶ Initialize parameters $\theta_j = \mathbf{w}_j, b_j$ randomly (note: not **0**)
- ▶ At each time step t :

Backpropagation

- ▶ For each (layer of) neuron(s) j :
 - ▶ Initialize parameters $\theta_j = \mathbf{w}_j, b_j$ randomly (note: not $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ For each (layer of) neuron(s) j , starting from the output and working backwards:
 - ▶ Compute gradient $\nabla_j L$

Backpropagation

- ▶ For each (layer of) neuron(s) j :
 - ▶ Initialize parameters $\theta_j = \mathbf{w}_j, b_j$ randomly (note: not $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ For each (layer of) neuron(s) j , starting from the output and working backwards:
 - ▶ Compute gradient $\nabla_j L$
 - ▶ For each (layer of) neuron(s) j :
 - ▶ Move in direction of negative gradient

Backpropagation

- ▶ For each (layer of) neuron(s) j :
 - ▶ Initialize parameters $\theta_j = \mathbf{w}_j, b_j$ randomly (note: not $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ For each (layer of) neuron(s) j , starting from the output and working backwards:
 - ▶ Compute gradient $\nabla_j L$
 - ▶ For each (layer of) neuron(s) j :
 - ▶ Move in direction of negative gradient
- ▶ Because L is not necessarily convex anymore, we are not guaranteed to reach a global minimum

Backpropagation

- ▶ For each (layer of) neuron(s) j :
 - ▶ Initialize parameters $\theta_j = \mathbf{w}_j, b_j$ randomly (note: not $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ For each (layer of) neuron(s) j , starting from the output and working backwards:
 - ▶ Compute gradient $\nabla_j L$
 - ▶ For each (layer of) neuron(s) j :
 - ▶ Move in direction of negative gradient
- ▶ Because L is not necessarily convex anymore, we are not guaranteed to reach a global minimum
 - ▶ But it works well enough in practice

Example

- ▶ Python time (again)!