

Contextualized Word Embeddings

CS114 Lab 12

Kenneth Lai

April 24, 2020

Contextualized Word Embeddings

CS114 Lab 12

Kenneth Lai

April 24, 2020



Source 1



Source 2

Contextualized Word Embeddings

CS114 Lab 12

Kenneth Lai

April 24, 2020



Source 1



Source 3

Word Embeddings

- ▶ Distributed representations of words

Word Embeddings

- ▶ Distributed representations of words
 - ▶ What is the difference between **distributed** and **distributional** representations?

Word Embeddings

- ▶ Sparse vectors

Word Embeddings

- ▶ Sparse vectors
 - ▶ One-hot, tf-idf, PPMI, etc.

Word Embeddings

- ▶ Sparse vectors
 - ▶ One-hot, tf-idf, PPMI, etc.
- ▶ Dense vectors

Word Embeddings


- ▶ Sparse vectors
 - ▶ One-hot, tf-idf, PPMI, etc.
- ▶ Dense vectors
 - ▶ SVD, word2vec, etc.

word2vec


- ▶ Based on a feedforward neural network language model

word2vec

- Based on a feedforward neural network language model




CBOW




Skip-gram

- Based on a feedforward neural network language model




CBOW




Skip-gram

- CBOW: use context to predict current word

- Based on a feedforward neural network language model



CBOW



Skip-gram

- CBOW: use context to predict current word
- Skip-gram: use current word to predict context

word2vec

- ▶ Input layer: one-hot word vectors

word2vec

- ▶ Input layer: one-hot word vectors
- ▶ Hidden (projection) layer: identity (linear!) activation function, no bias

word2vec

- ▶ Input layer: one-hot word vectors
- ▶ Hidden (projection) layer: identity (linear!) activation function, no bias
 - ▶ Input \rightarrow hidden = table lookup (in weight matrix)

word2vec

- ▶ Input layer: one-hot word vectors
- ▶ Hidden (projection) layer: identity (linear!) activation function, no bias
 - ▶ Input \rightarrow hidden = table lookup (in weight matrix)
- ▶ Output layer: softmax activation function

word2vec

- ▶ Input layer: one-hot word vectors
- ▶ Hidden (projection) layer: identity (linear!) activation function, no bias
 - ▶ Input \rightarrow hidden = table lookup (in weight matrix)
- ▶ Output layer: softmax activation function
- ▶ Word embeddings: output of hidden layer (= hidden layer weights corresponding to input word)

word2vec

- ▶ Input layer: one-hot word vectors
- ▶ Hidden (projection) layer: identity (linear!) activation function, no bias
 - ▶ Input \rightarrow hidden = table lookup (in weight matrix)
- ▶ Output layer: softmax activation function
- ▶ Word embeddings: output of hidden layer (= hidden layer weights corresponding to input word)

word embedding



Polysemy

- ▶ One vector per word type

Polysemy

- ▶ One vector per word type
- ▶ But words have multiple senses

Polysemy

- ▶ One vector per word type
- ▶ But words have multiple senses
 - ▶ a **mouse**¹ controlling a computer system in 1968

Polysemy

- ▶ One vector per word type
- ▶ But words have multiple senses
 - ▶ a **mouse**¹ controlling a computer system in 1968
 - ▶ a quiet animal like a **mouse**²

Polysemy

- ▶ One vector per word type
- ▶ But words have multiple senses
 - ▶ a **mouse**¹ controlling a computer system in 1968
 - ▶ a quiet animal like a **mouse**²
- ▶ Should **mouse**¹ and **mouse**² have the same word embedding?

Polysemy

- ▶ One vector per word type
- ▶ But words have multiple senses
 - ▶ a **mouse**¹ controlling a computer system in 1968
 - ▶ a quiet animal like a **mouse**²
- ▶ Should **mouse**¹ and **mouse**² have the same word embedding?
 - ▶ Why not?

Polysemy

- ▶ One vector per word type
- ▶ But words have multiple senses
 - ▶ ... mouse¹ ... computer ...
 - ▶ ... animal ... mouse² ...
- ▶ Should mouse¹ and mouse² have the same word embedding?
 - ▶ Why not?
 - ▶ Syntagmatic association between mouse¹ and computer

Polysemy

- ▶ One vector per word type
- ▶ But words have multiple senses
 - ▶ ... mouse¹ ... computer ...
 - ▶ ... animal ... mouse² ...
- ▶ Should mouse¹ and mouse² have the same word embedding?
 - ▶ Why not?
 - ▶ Syntagmatic association between mouse¹ and computer
 - ▶ Syntagmatic association between mouse² and animal

Polysemy

- ▶ One vector per word type
- ▶ But words have multiple senses
 - ▶ ... mouse¹ ... computer ...
 - ▶ ... animal ... mouse² ...
- ▶ Should mouse¹ and mouse² have the same word embedding?
 - ▶ Why not?
 - ▶ Syntagmatic association between mouse¹ and computer
 - ▶ Syntagmatic association between mouse² and animal
 - ▶ Paradigmatic association between computer and animal!


Polysemy

- ▶ One vector per word type
- ▶ But words have multiple senses
 - ▶ ... mouse¹ ... computer ...
 - ▶ ... animal ... mouse² ...
- ▶ Should mouse¹ and mouse² have the same word embedding?
 - ▶ Why not?
 - ▶ Syntagmatic association between mouse¹ and computer
 - ▶ Syntagmatic association between mouse² and animal
 - ▶ Paradigmatic association between computer and animal!
- ▶ How can we distinguish between mouse¹ and mouse²?

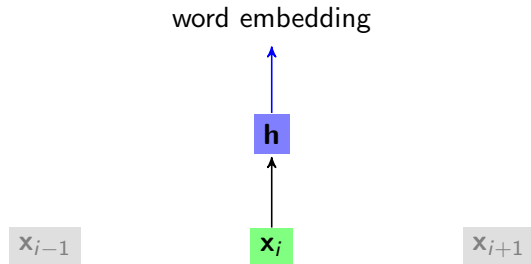
Polysemy

- ▶ One vector per word type
- ▶ But words have multiple senses
 - ▶ ... mouse¹ ... computer ...
 - ▶ ... animal ... mouse² ...
- ▶ Should mouse¹ and mouse² have the same word embedding?
 - ▶ Why not?
 - ▶ Syntagmatic association between mouse¹ and computer
 - ▶ Syntagmatic association between mouse² and animal
 - ▶ Paradigmatic association between computer and animal!
- ▶ How can we distinguish between mouse¹ and mouse²?
 - ▶ Context!

Word Embeddings




Word Embeddings



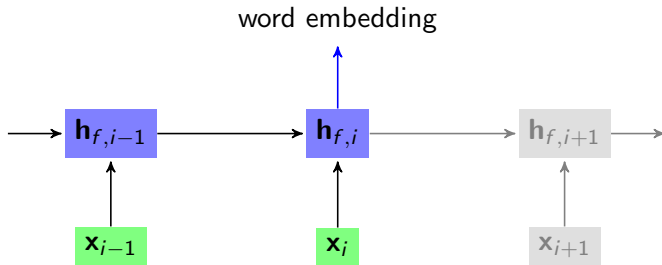
- h is an embedding of x_i only

Word Embeddings



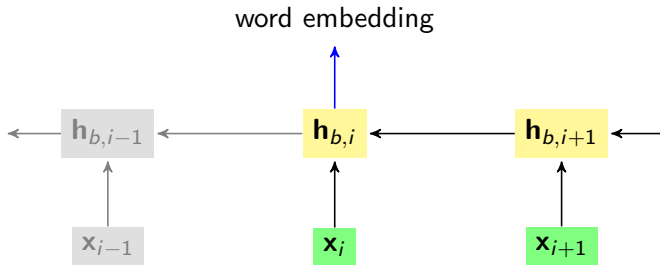
- ▶ h is an embedding of x_i only
 - ▶ How can we embed context information in h ?

Word Embeddings



- ▶ **h** is an embedding of x_i only
 - ▶ How can we embed context information in **h**?

Word Embeddings




- ▶ \mathbf{h} is an embedding of \mathbf{x}_i only
 - ▶ How can we embed context information in \mathbf{h} ?



► Embeddings from Language Models



- ▶ Embeddings from Language Models
- ▶ Based on a bidirectional recurrent neural network language model





- ▶ Input layer: pre-trained word vectors (e.g. from word2vec)



- ▶ Input layer: pre-trained word vectors (e.g. from word2vec)
- ▶ 2 bidirectional LSTM layers




- ▶ Input layer: pre-trained word vectors (e.g. from word2vec)
- ▶ 2 bidirectional LSTM layers
- ▶ Output layer: softmax



- ▶ Input layer: pre-trained word vectors (e.g. from word2vec)
- ▶ 2 bidirectional LSTM layers
- ▶ Output layer: softmax

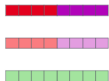
- ▶ Word embeddings: weighted sum of outputs of input and LSTM layers (task dependent)





Embedding of “stick” in “Let’s stick to” - Step #2

1- Concatenate hidden layers



2- Multiply each vector by a weight based on the task




3- Sum the (now weighted) vectors




ELMo embedding of “stick” for this task in this context

Forward Language Model



Backward Language Model




Source



► Bidirectional Encoder Representations from Transformers



► Bidirectional Encoder Representations from Transformers



Transformers

- ▶ An encoder-decoder architecture with attention

Transformers

- ▶ An **encoder-decoder architecture** with **attention**
- ▶ To learn more about encoder-decoder architecture, take StatNLP in the fall!

Transformers

- ▶ An **encoder-decoder architecture** with **attention**
- ▶ To learn more about encoder-decoder architecture, take StatNLP in the fall!
- ▶ Encoders take a sequence as input, output sequence representation

Transformers

- ▶ An **encoder-decoder architecture** with **attention**
- ▶ To learn more about encoder-decoder architecture, take StatNLP in the fall!
- ▶ Encoders take a sequence as input, output sequence representation
- ▶ Decoders take a sequence representation as input, output sequence

Transformers

- ▶ An **encoder-decoder architecture** with **attention**
- ▶ To learn more about encoder-decoder architecture, take StatNLP in the fall!
- ▶ Encoders take a sequence as input, output sequence representation
- ▶ Decoders take a sequence representation as input, output sequence
- ▶ Useful for machine translation, among other things

Transformers

- ▶ An **encoder-decoder architecture** with **attention**
- ▶ To learn more about encoder-decoder architecture, take StatNLP in the fall!
- ▶ Encoders take a sequence as input, output sequence representation
- ▶ Decoders take a sequence representation as input, output sequence
- ▶ Useful for machine translation, among other things
- ▶ BERT only uses the encoder part

Attention

- ▶ To learn more about attention, take StatNLP in the fall!

Attention

- ▶ To learn more about attention, take StatNLP in the fall!
- ▶ For each input i , compute a **query**, **key**, and **value**

Attention

- ▶ To learn more about attention, take StatNLP in the fall!
- ▶ For each input i , compute a **query**, **key**, and **value**
- ▶ The output for each input i is a weighted sum of the values for all inputs j , by how similar the key for j is to the query for i

Attention

- ▶ To learn more about attention, take StatNLP in the fall!
- ▶ For each input i , compute a **query**, **key**, and **value**
- ▶ The output for each input i is a weighted sum of the values for all inputs j , by how similar the key for j is to the query for i
- ▶ Result: the output for each input i contains information about the whole sequence

Attention

- ▶ To learn more about attention, take StatNLP in the fall!
- ▶ For each input i , compute a **query**, **key**, and **value**
- ▶ The output for each input i is a weighted sum of the values for all inputs j , by how similar the key for j is to the query for i
- ▶ Result: the output for each input i contains information about the whole sequence
 - ▶ More information about more relevant parts of the sequence

Attention

- ▶ To learn more about attention, take StatNLP in the fall!
- ▶ For each input i , compute a **query**, **key**, and **value**
- ▶ The output for each input i is a weighted sum of the values for all inputs j , by how similar the key for j is to the query for i
- ▶ Result: the output for each input i contains information about the whole sequence
 - ▶ More information about more relevant parts of the sequence
- ▶ Useful for long-distance dependencies, among other things



- ▶ Input layer: pre-trained word vectors (e.g. from word2vec)




- ▶ Input layer: pre-trained word vectors (e.g. from word2vec)
- ▶ 12-24 encoder layers



- ▶ Input layer: pre-trained word vectors (e.g. from word2vec)
- ▶ 12-24 encoder layers
 - ▶ Encoder layer = (shared) attention layer + (individual) feedforward layers




- ▶ Input layer: pre-trained word vectors (e.g. from word2vec)
- ▶ 12-24 encoder layers
 - ▶ Encoder layer = (shared) attention layer + (individual) feedforward layers






- ▶ Input layer: pre-trained word vectors (e.g. from word2vec)
- ▶ 12-24 encoder layers
 - ▶ Encoder layer = (shared) attention layer + (individual) feedforward layers



- ▶ Output layer: 2 pre-training tasks




- ▶ Input layer: pre-trained word vectors (e.g. from word2vec)
- ▶ 12-24 encoder layers
 - ▶ Encoder layer = (shared) attention layer + (individual) feedforward layers



- ▶ Output layer: 2 pre-training tasks
 - ▶ Masked LM (Cloze)



- ▶ Input layer: pre-trained word vectors (e.g. from word2vec)
- ▶ 12-24 encoder layers
 - ▶ Encoder layer = (shared) attention layer + (individual) feedforward layers



- ▶ Output layer: 2 pre-training tasks
 - ▶ Masked LM (Cloze)
 - ▶ NSP (Next Sentence Prediction)

Masked LM


- ▶ Mask 15% of input tokens at random

Masked LM

- ▶ Mask 15% of input tokens at random
- ▶ Predict masked words

Masked LM

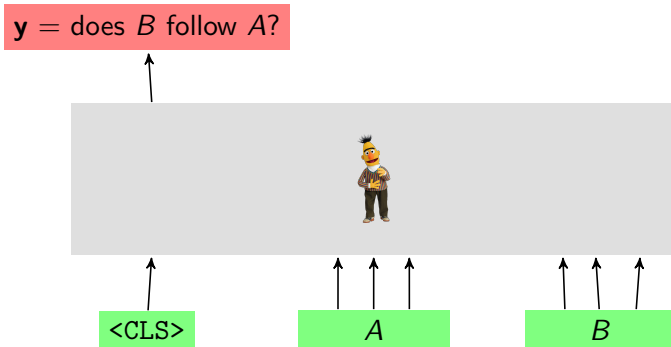
- ▶ Mask 15% of input tokens at random
- ▶ Predict masked words



- ▶ Given sentences A and B , does B follow A ?

NSP

- ▶ Given sentences A and B , does B follow A ?




























- ▶ Word embeddings: combinations of outputs of encoder layers



► Word embeddings: combinations of outputs of encoder layers

What is the best contextualized embedding for “Help” in that context?

For named-entity recognition task CoNLL-2003 NER

		Dev F1 Score
12 	First Layer Embedding 	91.0
• • •	Last Hidden Layer 12 	94.9
7 	12 	
6 	+ • • •	
5 	+ 2 	
4 	+ 1 	95.5
3 	= 	
2 	Second-to-Last Hidden Layer 11 	95.6
1 		
	Sum Last Four Hidden 12 	
Help	+ 11 	
	+ 10 	
	+ 9 	95.9
	= 	
	Concat Last Four Hidden  9 10 11 12 	96.1